

## Homework 3

### Due Sunday, May 1st via GradeScope

---

**Problem 1: Recurrence** Implement a simple CUDA program for a recurrence relation (inspired by the Mandelbrot Set) for many different starting points.

- **1.1 Allocate GPU memory** Idea: Use `cudaMalloc` to allocate memory on the GPU device for both the input and output arrays we will need for the recurrence implementation. Free memory with `cudaFree` at end of `main()`.

---

```
1      // Allocate num_bytes of memory to the device arrays
2      cudaMalloc(&device_input_array, num_bytes);
3      cudaMalloc(&device_output_array, num_bytes);
4      ...
5      ...
6
7      // Deallocate memory from both device arrays
8      cudaFree(device_input_array);
9      cudaFree(device_output_array);
```

---

- **1.2 Initialize array of random floats** Idea: Use in-built `rand()` functor from the standard `cpp` library and scale it between -1 and 1 as required.

---

```
1      // Initialize an array of size arr_size in input_array with
2      // random floats between -1 and 1
3      void initialize_array(vec &input_array, size_t arr_size) {
4          input_array.resize(arr_size);
5          std::generate(input_array.begin(), input_array.end(), rand);
6
7          for(int i=0; i<arr_size; i++){
8              input_array[i] = static_cast<float>(input_array[i])/(RAND_MAX/2)-1;
9          }
10     }
```

---

- **1.3 Implement recurrence kernel** Idea: Recurrence operations themselves are not independent and therefore not parallelizable, however, we can parallelize doing many of these recurrence loops for different starting points (constants). So, we want to parallelize over our 1-dim input array of constants.

Since our kernel needs to handle cases where the number of threads is less than the number of entries in our input array, we need to use a grid-stride loop.

---

```
1      /**
2      * Implement the kernel recurrence.
3      * The CPU implementation is in host_recurrence() in main_q1.cu.
```

```

4      */
5      __global__ void recurrence(const elem_type* input_array,
6                                elem_type* output_array,
7                                size_t num_iter,
8                                size_t array_length) {
9
10         for (int xid = blockIdx.x * blockDim.x + threadIdx.x;
11              xid < array_length;
12              xid += blockDim.x * gridDim.x) {
13
14             elem_type z = 0;
15             elem_type constant = input_array[xid];
16
17             int it=0;
18             while(it<num_iter) {
19                 z = z * z + constant;
20                 it++;
21             }
22             output_array[xid] = z;
23         }
24     }

```

Console logs.

Starting at Fri Apr 29 00:28:45 UTC 2022

nvcc -O3 -std=c++11 -arch=compute\_75 -code=sm\_75 -o main\_q1 main\_q1.cu

Output from main\_q1

-----

Largest error found at pos: 15 error 7.81565e-08 expected 1.52526 and got 1.525  
 Largest error found at pos: 0 error 0 expected 0.680375 and got 0.680375  
 Largest error found at pos: 439038 error 1.19193e-07 expected 1.00014 and got 1  
 Largest error found at pos: 142710 error 2.38333e-07 expected 2.00072 and got 2  
 Largest error found at pos: 482709 error 5.61004e-07 expected 16.9994 and got 1  
 Largest error found at pos: 482709 error 1.15797e-06 expected 289.897 and got 2  
 Largest error found at pos: 482709 error 2.324e-06 expected 84041.4 and got 840  
 Largest error found at pos: 482709 error 4.63941e-06 expected 7.06296e+09 and g  
 Largest error found at pos: 482709 error 9.25711e-06 expected 4.98854e+19 and g  
 Largest error found at pos: 138972 error 1.79297e-05 expected 8.03779e+21 and g  
 Largest error found at pos: 905817 error 2.59306e-05 expected 1.66519e+35 and g

Questions 1.1-1.3: your code passed all the tests!

Q1.4

-----

Number of Threads	Performance TFlops/sec
32	0.843698

64	1.71319
96	2.5743
128	3.42082
160	4.4214
192	5.79964
224	6.8216
256	7.70162
288	8.27133
320	9.18469
352	10.0615
384	10.7552
416	10.529
448	11.2155
480	11.9355
512	12.3663
544	11.3325
576	11.5058
608	12.4919
640	13.0052
672	11.5092
704	12.075
736	12.695
768	12.7192
800	11.7815
832	12.4544
864	12.4641
896	13.1923
928	12.4277
960	12.2472
992	12.8563
1024	13.0942

Q1.5

Number of Blocks	Performance TFlops/sec
36	1.68669
72	3.35125
108	4.43883
144	5.86105
180	6.8456
216	8.1191
252	8.12224
288	9.39581
324	7.71112
360	8.63641

396	8.93084
432	9.58119
468	9.36923
504	9.9853
540	9.35853
576	9.94791
612	9.56487
648	9.41304
684	9.93957
720	10.5008
756	10.0355
792	10.6237
828	10.1014
864	11.9466
900	12.516
936	12.909
972	12.8201
1008	12.9218
1044	12.5967
1080	12.8112
1116	12.4403
1152	12.7656

## Q1.6

---

Number of Iters	Performance TFlops/sec
20	1.54131
40	3.29381
60	4.80154
80	5.64972
100	6.20655
120	7.32422
140	5.63426
160	8.15661
180	8.3581
200	8.83392
300	6.81323
400	8.40619
500	8.00461
600	8.20031
700	8.46228
800	8.77809
900	9.48087
1000	9.32557
1200	9.32836

1400	9.49436
1600	9.68148
1800	9.66993
2000	9.91434
2200	9.97533
2400	10.0087
2600	10.1537
2800	10.083
3000	10.2325

- **1.4 Vary number of threads per block ...** Many threads concurrently trying to access cache (traffic)  
[Insert graph]
- **1.5 Vary number of blocks**
- **1.6 Vary number of iterations**

## Problem 2: PageRank ...

- **2.1: ...**

Submission information logs.

```
jelc@cardinal1:~$ /afs/ir.stanford.edu/class/cme213/script/submit.py hw2 private/cme213-
Submission for assignment 'hw2' as user 'jelc'
```

```
Attempt 1/10
```

```
Time stamp: 2022-04-13 20:09
```

```
List of files being copied:
```

```
private/cme213-jelc53/hw2/sum.h [768 bytes]
private/cme213-jelc53/hw2/parallel_radix_sort.h [7625 bytes]
```

```
Your files were copied successfully.
```

```
Directory where files were copied: /afs/ir.stanford.edu/class/cme213/submissions/hw2/jelc/1
```

```
List of files in this directory:
```

```
sum.h [768 bytes]
parallel_radix_sort.h [7625 bytes]
metadata [137 bytes]
```

This completes the submission process. Thank you!

```
jelc@cardinal1:~$ ls /afs/ir.stanford.edu/class/cme213/submissions/hw2/jelc/1
metadata parallel_radix_sort.h sum.h
```