

Homework 4: CUDA GPU Matrix Operations

Goal of the homework is to implement a finite difference solver for the 2-dim heat equation using CUDA GPU programming.

Problem 1: Implement gpuStencil Global Idea: parallelize spatial dimension updates for each time step iteration. One key challenge was to handle the different border sizes for order inputs of 2, 4 and 8. Kernel logic included below.

```
1  /**
2   * Kernel to propagate finite difference grid from the current
3   * time point to the next.
4   *
5   * This kernel should be very simple and only use global memory
6   * and 1d threads and blocks.
7   *
8   * @param next[out] Next grid state.
9   * @param curr Current grid state.
10  * @param gx Number of grid points in the x dimension.
11  * @param nx Number of grid points in the x dimension to which the full
12  *           stencil can be applied (ie the number of points that are at least
13  *           order/2 grid points away from the boundary).
14  * @param ny Number of grid points in the y dimension to which the full
15  *           stencil can be applied.
16  * @param xcfl Courant number for x dimension.
17  * @param ycfl Courant number for y dimension.
18  */
19  template<int order>
20  __global__
21  void gpuStencilGlobal(float* next, const float* __restrict__ curr,
22                        int gx, int nx, int ny, float xcfl, float ycfl) {
23
24      int borderSize = (int) (order / 2);
25      int i = blockIdx.x * blockDim.x + threadIdx.x;
26
27      if (i < nx*ny) {
28          int x = borderSize + (int) (i / nx);
29          int y = borderSize + (i % nx);
30          int idx = gx * y + x;
31          next[idx] = Stencil<order>(&curr[idx], gx, xcfl, ycfl);
32      }
33  }
34 }
```

3D surface plots of temperature on 256 x 256 grid at iterations 0, 1000 and 2000 respectively, with 8th order. To do this, I used parameter settings of: order = 8 and nx = ny = 248.

```
1  // Allocate num_bytes of memory to the device arrays
```

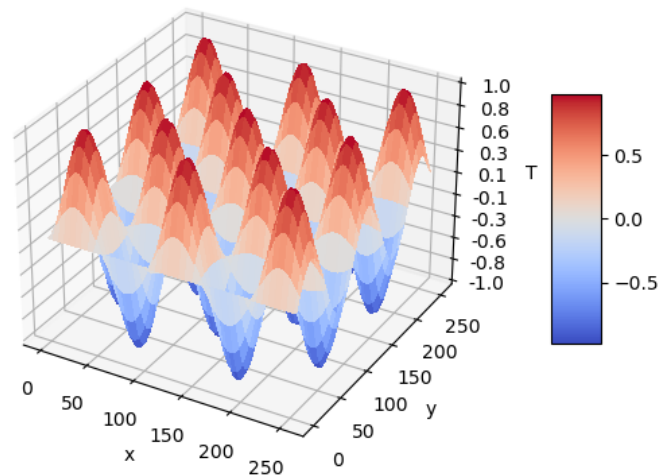


Figure 1: 3D surface plot of temperature at iteration 0

```

2  cudaMalloc(&device_input_array, num_bytes);
3  cudaMalloc(&device_output_array, num_bytes);
4  ...
5  ...
6
7  // Deallocate memory from both device arrays
8  cudaFree(device_input_array);
9  cudaFree(device_output_array);

```

Submission information logs.

```
jelc@cardinal2:~$ /afs/ir.stanford.edu/class/cme213/script/submit.py hw3 private/cme213-
Submission for assignment 'hw3' as user 'jelc'
```

Attempt 2/10

Time stamp: 2022-05-01 21:36

List of files being copied:

```

private/cme213-jelc53/hw3/main_q1.cu  [13253 bytes]
private/cme213-jelc53/hw3/recurrence.cuh  [1589 bytes]
private/cme213-jelc53/hw3/pagerank.cuh  [5894 bytes]
private/cme213-jelc53/hw3/benchmark.cuh  [795 bytes]

```

Your files were copied successfully.

Directory where files were copied: /afs/ir.stanford.edu/class/cme213/submissions/hw3/jelc

List of files in this directory:

```

main_q1.cu  [13253 bytes]
recurrence.cuh  [1589 bytes]

```

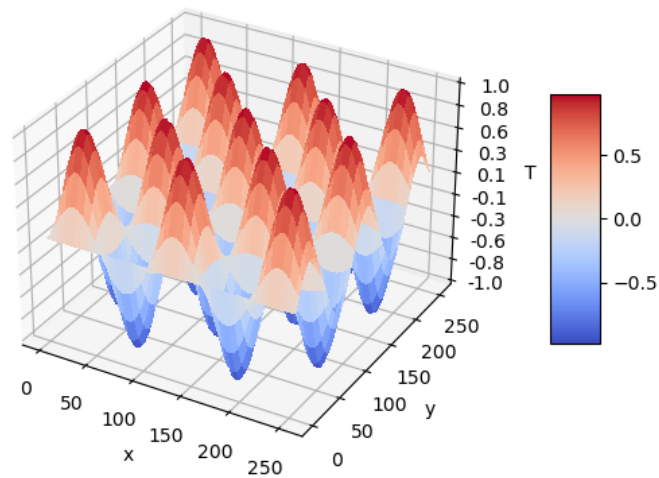


Figure 2: 3D surface plot of temperature at iteration 1000

```
pagerank.cuh [5894 bytes]
benchmark.cuh [795 bytes]
metadata [137 bytes]
```

This completes the submission process. Thank you!

```
jelc@cardinal2:~$ ls /afs/ir.stanford.edu/class/cme213/submissions/hw3/jelc/2
benchmark.cuh  main_q1.cu  metadata  pagerank.cuh  recurrence.cuh
```

”